

I/O Performance trade-offs among RNTuple's persistent layouts for DUNE Data Products

S M Shovan, *FCSI Summer Intern '25*

July 31, 2025

Fermi National Accelerator Laboratory

Introduction

Persistent Layouts

Parallel Optimizations

Challenges

Results

Trade-offs

Introduction

Introduction

- The Deep Underground Neutrino Experiment (DUNE) is projected to record roughly 30 PB of liquid-argon TPC data per year [1]—far beyond the scale of previous neutrino experiments.
- DUNE's **Phlex** test-stand at Fermilab provides realistic hit/wire samples to prototype read-out and offline storage.
- ROOT's new `RNTuple` backend is a candidate for the long-term DUNE data model, promising faster compression, cluster-aware reads, and thread-safe writes.
- This study benchmarks alternative `RNTuple` *persistent layouts* (AOS/SOA, vertical splits, granularity levels) using Phlex-generated Hit and Wire data products.
- Focus data products: `recob::Hit` (charge deposits) and `recob::Wire` (ROI-compressed waveforms).

- A single DUNE far-detector module streams about 1.2 TB/s of raw data before compression [2]; naive storage could potentially overwhelm the archival budget.
- Efficient layout choice can cut file size and accelerate cluster reads needed for GPU/CPU reconstruction farms.

Problem Statement

- Which RNTuple persistent layout minimises *both* write time and on-disk footprint for DUNE Hit/Wire hierarchies?
- How does vertical splitting (all-DP, per-DP, per-group) interact with horizontal granularities (event, spill, element) under Phlex workloads?

Objectives

- Benchmark seven layout variants on a 1 M-event Phlex dataset (`recob::Hit`, `recob::Wire` with ROIs).
- Measure: write throughput, cold/warm read latency, compressed file size, and multi-thread scaling (1–64 threads).
- Quantify trade-offs of ROI flattening, vertical split depth, and SOA vs. AOS in the context of DUNE's access patterns.

Persistent Layouts

Data Layouts: AOS vs. SOA

Array of Structures (AOS)

Stores complete objects in an array.

```
struct Hit {  
    long long EventID;  
    unsigned int fChannel;  
    float fPeakTime;  
};  
// Array: [Hit1, Hit2, ...]
```

Example: Hit, Wire (per-item entries).

Structure of Arrays (SOA)

Separate arrays per field.

```
struct Hits {  
    vector<unsigned int> fChannel;  
    vector<float> fPeakTime;  
};  
// Columns: fChannel[ ],  
fPeakTime[ ], ...
```

Example: Hits, Wires (per-event vectors).

Layout Strategies




Figure 1: AOS Layout

Layout Strategies

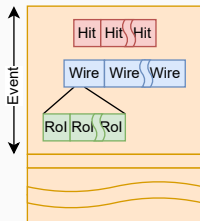


Figure 1: AOS Layout

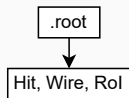


Figure 2: 1 RNTuple for all data products

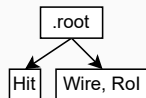


Figure 3: 1 RNTuple per data product

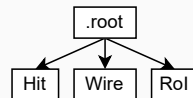


Figure 4: 1 RNTuple per group

Layout Strategies

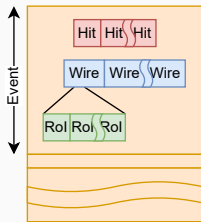


Figure 1: AOS Layout

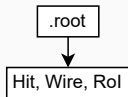


Figure 2: 1 RNTuple for all data products

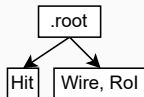


Figure 3: 1 RNTuple per data product

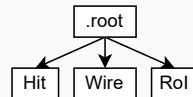


Figure 4: 1 RNTuple per group




Figure 5: 1 fill/row per event

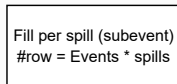


Figure 6: 1 fill/row per spill




Figure 7: 1 fill/row per top object

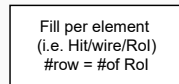


Figure 8: 1 fill/row per element

Layout Strategies: AOS vs SOA

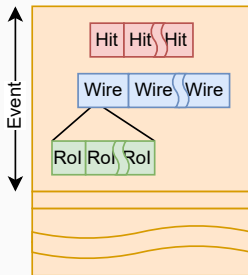


Figure 9: AOS Layout




Figure 10: SOA Layout

Granularity vs. Vertical Split Matrix

Horizontal Granularity	1 NTuple (all DP)	1 NTuple / DP	1 NTuple / group
Event-wise	<code>event_allDP()</code>	<code>event_perDP()</code>	<code>event_perGroup()</code>
Spill-wise	<code>spill_allDP()</code>	<code>spill_perDP()</code>	<code>spill_perGroup()</code>
Top-object-wise	–	<code>topObject_perDP()</code>	<code>topObject_perGroup()</code>
Element-wise	–	<code>element_perDP()</code>	<code>element_perGroup()</code>

DP = Data Product

Memory/Data Considerations

Memory/Data Considerations

Parallel Optimizations

Write Optimization: Multi-Threaded Chunking

Parallel Chunking Divide events into thread-specific ranges for concurrent filling.

```
std::vector<unsigned int> seeds =  
generateSeeds(nThreads);  
for (int th = 0; th < nThreads; ++th) {  
    int first = th * chunkSize;  
    int last = std::min(first + chunkSize,  
totalEvents);  
    futures.emplace_back(std::async(  
std::launch::async, thinWorkFunc, first,  
last, seeds[th], th ));}
```

Example: `executeInParallel` writers.

Synchronized Filling Use mutex to safely flush clusters after filling.

```
{ std::lock_guard<std::mutex> lock(mutex);  
}  
hitContext.FlushCluster();
```

Project Use: Scales writes with cores for large datasets.

Read Optimization: Cluster-Aware Splitting

Cluster Splitting Split read ranges by cluster boundaries to avoid duplicates.

```
auto clusters =  
split_range_by_clusters(*reader, nChunks);  
for (auto& chunk : clusters) {  
    futures.push_back(std::async(  
        &processChunk, chunk.first, chunk.second  
    ));  
}
```

Benefits Prevents redundant reads, optimizes multi-threaded access.

```
std::vector<std::pair<size_t, size_t>>  
split_range_by_clusters(ROOT::RNTupleReader&  
reader, int nChunks)
```

Project Use: Enhances read efficiency.

Challenges

Parallel Write Challenge: File Corruption Solution

Problem: Concurrent Flushes

Unsynchronized cluster flushes cause file corruption in multi-threaded writes. **Example Issue:** Threads overwriting shared file regions.

Solution: Mutex Synchronization Lock during flushes to serialize access per cluster.

```
for (int idx = first; idx < last; ++idx) {  
    // Generate data for hits/wires  
    if (hitStatus.ShouldFlushCluster()) {  
        hitContext.FlushColumns();  
        { std::lock_guard<std::mutex> lock(mutex); }  
        hitContext.FlushCluster();  
    }  
}
```

Project Use: Ensures thread-safe parallel writes without corruption.

ROI Flattening vs. Custom Dictionary

ROI Flattening (Non-Dictionary)

Flattens hierarchical ROI data into vectors for efficient storage without custom classes.

```
struct Wires {  
    vector<unsigned int> fSignalROI_nROIs;  
    vector<size_t> fSignalROI_offsets;  
    vector<float> fSignalROI_data;  
};
```

Example: Used in non-dictionary experiments for raw vector-based I/O.

Custom Dictionary (ROOT Classes)

Uses structured classes with ClassDef for ROOT's dictionary system, enabling object-oriented I/O.

```
struct RegionOfInterest {  
    size_t offset;  
    vector<float> data;  
    ClassDef(RegionOfInterest, 3)  
};
```

Example: Used in dictionary experiments for type-safe, hierarchical data handling.

Results

- **Write Throughput:** Total events per second during RNTuple serialization.
- **Cold Read Time:** Latency for first access after file creation, reflecting raw I/O.
- **Warm Read Time:** Latency when data is cached, measuring memory locality effects.
- **Compressed File Size:** Total on-disk footprint post-write, accounting for RNTuple's column-wise compression.
- **Multi-thread Scaling:** Performance variation from 1 to 64 threads for write and read workloads.

AOS vs SOA: Write Performance Comparison




Figure 11: Write Performance Comparison: AOS vs SOA approaches across different data organization strategies. Red bars represent AOS (Array of Structures) performance, while blue bars show SOA (Structure of Arrays) performance.

CPU Time Overheads

Writer	A	B	C
event_perGroup	✓	✓	
spill_perGroup	✓	✓	✓
element_perData	✓	✓	
topObject_perData		✓	
event_perData		✓	
spill_perData		✓	✓

- **A – ROI flattening:** dominant for *element* and *perGroup* writers.
- **B – Multiple NTuples:** extra Fill/Flush cycles & mutex contention.
- **C – Spill re-partitioning:** additional work unique to *spill* writers.

AOS vs SOA: Cold Read Performance Comparison

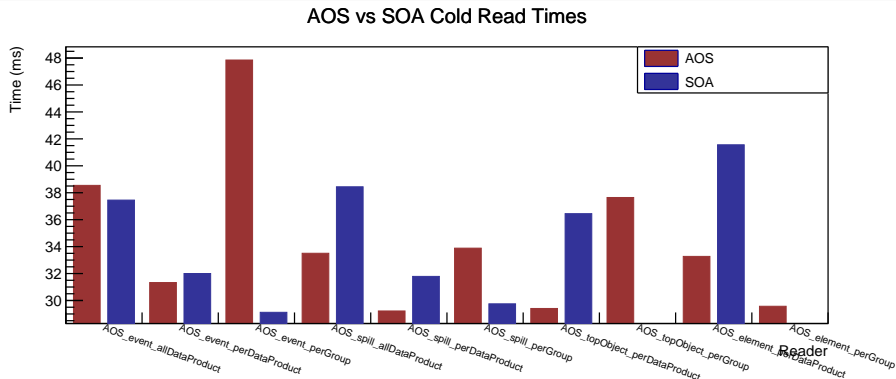


Figure 12: Cold Read Performance Comparison: Initial read times for AOS vs SOA implementations. Cold reads represent the first access to data after file creation, measuring baseline I/O performance.

AOS vs SOA: Warm Read Performance Comparison




Figure 13: Warm Read Performance Comparison: Subsequent read times for AOS vs SOA implementations. Warm reads benefit from cached data, showing optimized access patterns and memory locality effects.

AOS vs SOA: File Size Comparison

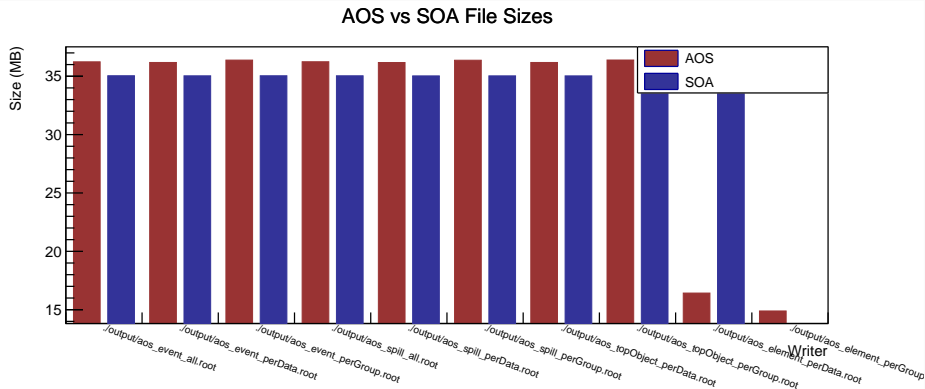


Figure 14: File Size Comparison: Storage efficiency of AOS vs SOA approaches. File sizes impact storage costs, network transfer times, and overall data management overhead in large scale HEP experiments.

Writer variant	Size (% of most)
aos_element_perData.root	47%
soa_element_perData.root	100%
aos_element_perGroup.root	42%
soa_element_perGroup.root	39%
aos_event_* / aos_spill_*	~103%
soa_event_* / soa_spill_*	100%

Key Take-aways

- Deep nesting & high row counts double basket overhead (offset tables, headers, CRC).
- Flattened perGroup writers cut size roughly in half by removing ROI nesting.

Trade-offs

Trade-off Analysis

Future Considerations

- [1] DUNE Collaboration, “Deep Underground Neutrino Experiment Technical Design Report–Volumell: DUNE Physics,” 2020, Sec.2.6.
- [2] DUNE Collaboration, “Data Acquisition System for the DUNE Far Detector,” IEEE NSS/MIC Proc., 2023.

Thank you!
Questions?