

# I/O Performance trade-offs among RNTuple's persistent layouts for DUNE Data Products

---

S M Shovan, *FCSI Summer Intern'25*

July 31, 2025

Fermi National Accelerator Laboratory

Introduction

Persistent Layouts

Parallel Optimizations

Challenges

Results

Trade-offs

# Introduction

---

# Introduction

- The Deep Underground Neutrino Experiment (DUNE) is projected to record roughly 30 PB of liquid-argon TPC data per year [1]—far beyond the scale of previous neutrino experiments.
- DUNE's **Phlex** test-stand at Fermilab provides realistic hit/wire samples to prototype read-out and offline storage.
- ROOT's new `RNTuple` backend is a candidate for the long-term DUNE data model, promising faster compression, cluster-aware reads, and thread-safe writes.
- This study benchmarks alternative `RNTuple` *persistent layouts* (AOS/SOA, vertical splits, granularity levels) for realistic data products.
- Focus data products: `recob::Hit` (charge deposits) and `recob::Wire` (ROI-compressed waveforms).

- A single DUNE far-detector module streams about 1.2 TB/s of raw data before compression [2]; naive storage could potentially overwhelm the archival budget.
- Efficient layout choice can cut file size and accelerate cluster reads needed for GPU/CPU reconstruction farms.

# Problem Statement

- Which RNTuple persistent layout minimises read time, write time and on-disk footprint for DUNE Hit/Wire hierarchies?
- How does vertical splitting such as one RNTuple for all data products interact with horizontal granularities (event, spill, element) under Phlex workloads?
- How does the choice of persistent layout affect the performance of the read and write operations?

# Objectives

- Benchmark seven layout variants on a 1 M-event Phlex dataset (`recob::Hit`, `recob::Wire` with ROIs).
- Measure: write throughput, cold/warm read latency, compressed file size, and multi-thread scaling (1–64 threads).
- Quantify trade-offs of ROI flattening, vertical split depth, and SOA vs. AOS.

## Persistent Layouts

---



# Data Layouts: AOS vs. SOA

## Array of Structures (AOS)

Stores complete objects in an array.

```
struct Hit {  
    long long EventID;  
    unsigned int fChannel;  
    float fPeakTime;  
};  
// Array: [Hit1, Hit2, ...]
```

**Example:** Hit, Wire (per-item entries).

## Structure of Arrays (SOA)

Separate arrays per field.

```
struct Hits {  
    vector<long long> EventID;  
    vector<unsigned int> fChannel;  
    vector<float> fPeakTime;  
};  
// Columns: EventID[ ], fChannel[ ], fPeakTime[ ], ...
```

**Example:** Hits, Wires (per-event vectors).

# Layout Strategies

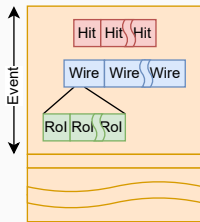


Figure 1: AOS Layout

# Layout Strategies

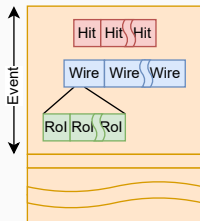


Figure 1: AOS Layout

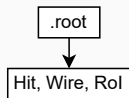


Figure 2: 1 RNTuple for all data products

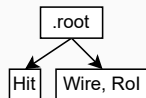


Figure 3: 1 RNTuple per data product

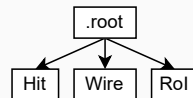


Figure 4: 1 RNTuple per group

# Layout Strategies

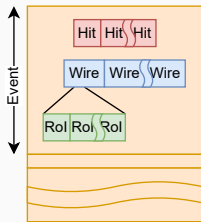


Figure 1: AOS Layout

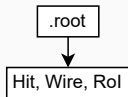


Figure 2: 1 RNTuple for all data products

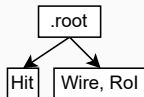


Figure 3: 1 RNTuple per data product

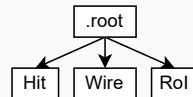


Figure 4: 1 RNTuple per group

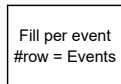


Figure 5: 1 fill/row per event

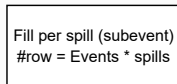


Figure 6: 1 fill/row per spill

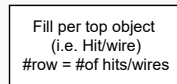


Figure 7: 1 fill/row per top object

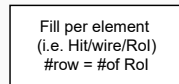


Figure 8: 1 fill/row per element

# Layout Strategies: AOS vs SOA

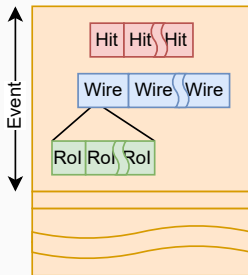


Figure 9: AOS Layout




Figure 10: SOA Layout

# Granularity vs. Vertical Split Matrix

Horizontal Granularity	1 NTuple (all DP)	1 NTuple / DP	1 NTuple / group
Event-wise	<code>event_allDP()</code>	<code>event_perDP()</code>	<code>event_perGroup()</code>
Spill-wise	<code>spill_allDP()</code>	<code>spill_perDP()</code>	<code>spill_perGroup()</code>
Top-object-wise	–	<code>topObject_perDP()</code>	<code>topObject_perGroup()</code>
Element-wise	–	<code>element_perDP()</code>	<code>element_perGroup()</code>

DP = Data Product

# Parallel Optimizations

---

# Write Optimization: Multi-Threaded Chunking

**Parallel Chunking** Divide events into thread-specific ranges for concurrent filling.

```
std::vector<unsigned int> seeds = generateSeeds(nThreads);  
for (int th = 0; th < nThreads; ++th) {  
    int first = th * chunkSize;  
    int last = std::min(first + chunkSize, totalEvents);  
    futures.emplace_back(std::async(  
        std::launch::async, thinWorkFunc, first, last, seeds[th], th ));  
}
```

**Example:** `executeInParallel` writers.

**Project Use:** Scales writes with cores for large datasets.



# Read Optimization: Cluster-Aware Splitting

**Cluster Splitting** Split read ranges by cluster boundaries to avoid duplicates.

```
auto clusters = split_range_by_clusters(*reader, nChunks);  
for (auto& chunk : clusters) {  
    futures.push_back(std::async(  
        &processChunk, chunk.first, chunk.second  
    ));  
}
```

**Helper Function** Defines cluster-based splits.

```
std::vector<std::pair<size_t, size_t>>  
split_range_by_clusters(ROOT::RNTupleReader& reader, int nChunks)
```

**Project Use:** Enhances read efficiency by reducing redundant reads.

## Challenges

---

# Parallel Write Challenge: File Corruption Solution

## Problem: Concurrent Flushes

Unsynchronized cluster flushes cause file corruption in multi-threaded writes. **Example Issue:** Threads overwriting shared file regions.

**Solution: Mutex Synchronization** Lock during flushes to serialize access per cluster.

```
for (int idx = first; idx < last; ++idx) {  
    // Generate data for hits/wires  
    if (hitStatus.ShouldFlushCluster()) {  
        hitContext.FlushColumns();  
        { std::lock_guard<std::mutex> lock(mutex); }  
        hitContext.FlushCluster();  
    }  
}
```

**Project Use:** Ensures thread-safe parallel writes without corruption.

# ROI Flattening vs. Custom Dictionary

## ROI Flattening (Non-Dictionary)

Flattens hierarchical ROI data into vectors for efficient storage without custom classes.

```
struct Wires {  
    vector<unsigned int> fSignalROI_nROIs;  
    vector<size_t> fSignalROI_offsets;  
    vector<float> fSignalROI_data;  
};
```

**Example:** Used in non-dictionary experiments for raw vector-based I/O.

## Custom Dictionary (ROOT Classes)

Uses structured classes with ClassDef for ROOT's dictionary system, enabling object-oriented I/O.

```
struct RegionOfInterest {  
    size_t offset;  
    vector<float> data;  
    ClassDef(RegionOfInterest, 3)  
};
```

**Example:** Used in dictionary experiments for type-safe, hierarchical data handling.

## Results

---

- **Write Throughput:** Total events per second during RNTuple serialization.
- **Cold Read Time:** Latency for first access after file creation, reflecting raw I/O.
- **Warm Read Time:** Latency when data is cached, measuring memory locality effects.
- **Compressed File Size:** Total on-disk footprint post-write, accounting for RNTuple's column-wise compression.
- **Multi-thread Scaling:** Performance variation from 1 to 64 threads for write and read workloads.

# AOS vs SOA: Write Performance Comparison




Figure 11: Write Performance Comparison: AOS vs SOA approaches across different data organization strategies. Red bars represent AOS (Array of Structures) performance, while blue bars show SOA (Structure of Arrays) performance.

# CPU Time Overheads

Writer	A	B	C
event_perGroup	✓	✓	
spill_perGroup	✓	✓	✓
element_perData	✓	✓	
topObject_perData		✓	
event_perData		✓	
spill_perData		✓	✓

- **A – ROI flattening:** dominant for \*element\* and \*perGroup\* writers.
- **B – Multiple NTuples:** extra Fill/Flush cycles & mutex contention.
- **C – Spill re-partitioning:** additional work unique to \*spill\* writers.



# AOS vs SOA: Cold Read Performance Comparison




Figure 12: Cold Read Performance Comparison: Initial read times for AOS vs SOA implementations.

# AOS vs SOA: Warm Read Performance Comparison




Figure 13: Warm Read Performance Comparison: Subsequent read times for AOS vs SOA implementations.

# AOS vs SOA: File Size Comparison

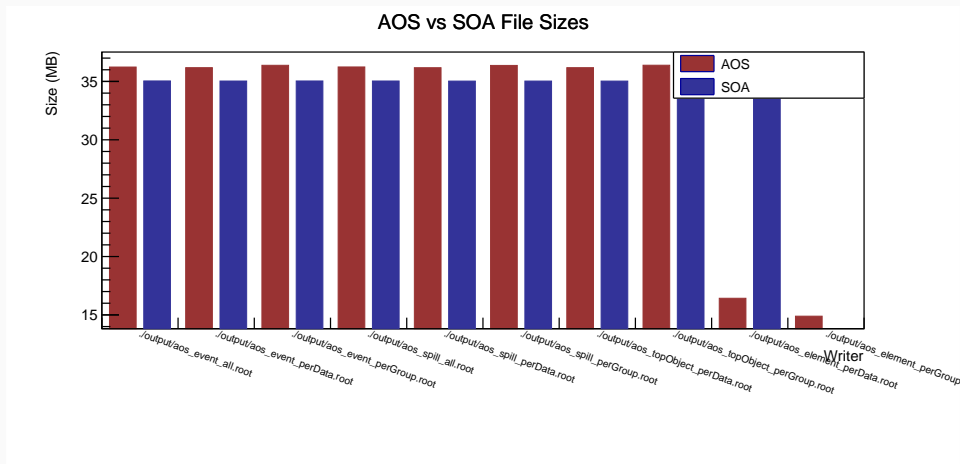


Figure 14: File Size Comparison: Storage efficiency of AOS vs SOA approaches.

# File Size Insights

Writer variant	Size (% of most)
aos_element_perData.root	47%
soa_element_perData.root	100%
aos_element_perGroup.root	42%
soa_element_perGroup.root	39%
aos_event_* / aos_spill_*	~103%
soa_event_* / soa_spill_*	100%

## Key Take-aways

- Deep nesting & high row counts double basket overhead (offset tables, headers, CRC).
- Flattened perGroup writers cut size roughly in half by removing ROI nesting.

## Trade-offs

---

# Trade-off Analysis

# Future Considerations

- [1] DUNE Collaboration, “Deep Underground Neutrino Experiment Technical Design Report–Volumell: DUNE Physics,” 2020, Sec.2.6.
- [2] DUNE Collaboration, “Data Acquisition System for the DUNE Far Detector,” IEEE NSS/MIC Proc., 2023.



**Thank you!**  
**Questions?**